

SC2001/CE2101/ CZ2101: Algorithm Design and Analysis


**Greedy Algorithms;
Dijkstra's Algorithm; Prim's Algorithm**

Instructor: Assoc. Prof. ZHANG Hanwang

Courtesy of Dr. Ke Yiping, Kelly's slides

Greedy Algorithms

Local optimality generally does **NOT** imply **global optimality**!!



Learning Objectives

At the end of this lecture, students should be able to:

- Explain the strategy of Greedy algorithms
- Solve single-source shortest paths problem using Dijkstra's algorithm
- Prove the correctness of Dijkstra's algorithm
- Describe Prim's algorithm for finding minimum spanning trees (MSTs)
- Prove the correctness of Prim's algorithm

Greedy Algorithms

- In optimization problems, the algorithm needs to make a series of choices whose overall effect is to minimize the total cost, or maximize the total benefit, of some system.
- There is a class of algorithms, called the **greedy algorithms**, in which we can find a solution by using only knowledge available at the time when the next choice (or guess) must be made.
- Each individual choice is the best within the knowledge available at the time.

Greedy Algorithms


- Each individual choice is not very expensive to compute.
- A choice cannot be undone, even if it is found to be a bad choice later.
- Greedy algorithms cannot guarantee to produce the optimal solution for a problem.

Dijkstra's Algorithm

Dijkstra's Algorithm

Shortest Path Problem:

The problem of finding the **shortest path** from **one vertex** in a graph G to **another vertex**. "Shortest" may be the least number of edges, or the least total weight, etc.



Dijkstra's Algorithm:

This is an algorithm to find the shortest paths from a single source vertex to all other vertices in a **weighted, directed** graph. All weights must be **nonnegative**.

Dijkstra's Algorithm

Dijkstra's algorithm keeps two sets of vertices:

- **S**: the set of vertices whose shortest paths from the source node have already been determined
 - they form the tree
- **V – S**: the remaining vertices

The other data structures needed are:

- **d**: an array of size $|V|$ to store the estimated lengths of shortest paths from the source node to all vertices
- **pi**: an array of size $|V|$ to store the predecessors for each vertex


Basic Steps

The basic steps are:


1. Initialise **d** and **pi**
2. Set **S** to empty
3. While there are still vertices in **V – S**
 - i. Move **u**, the vertex in **V – S** that has the shortest path estimate from source, to **S**
 - ii. For all the vertices in **V – S** that are connected to **u**, update their estimates of shortest distances to the source

A Toy Example

S: { }



Shortest paths from
s to other vertices



d	s	x	u	v	y


pi	s	x	u	v	y

Pseudocode of Dijkstra's Algorithm

```
Dijkstra_ShortestPath ( Graph G, Node source ) {  
  for each vertex v {  
    d[v] = infinity;  
    pi[v] = null pointer;  
    S[v] = 0;    // S[v] is 1 if v is in S  
  }             // S[v] is 0 if v is not in S  
  d[source] = 0;  
  put all vertices in priority queue, Q, in d[v]'s increasing order;  
  while not Empty(Q) {  
    u = ExtractCheapest(Q);  
    S[u] = 1;    /* Add u to S */  
  }  
}
```

Pseudocode of Dijkstra's Algorithm

```
for each vertex  $v$  adjacent to  $u$ 
  if ( $S[v] \neq 1$  and  $d[v] > d[u] + w[u, v]$ ) {
    remove  $v$  from  $Q$ ;
     $d[v] = d[u] + w[u, v]$ ;
     $pi[v] = u$ ;
    insert  $v$  into  $Q$  according to its  $d[v]$ ;
  }
} // end of while loop
}
```




Worst case time complexity of Dijkstra's algorithm is $O(|V|^2)$ (analysis not required).

Proof of Correctness

Property of Shortest Path

Lemma 1: In a weighted graph G , suppose that a shortest path from x to z consists of a path P from x to y followed by a path Q from y to z . Then P is a shortest path from x to y and Q is a shortest path from y to z .



Proof (By Contradiction):

Assume that P is not the shortest path from x to y . Then there will be another path from x to y , P' which is shorter than P . As a result P' followed by Q will be a path **shorter** than P followed by Q . But it was known that P followed by Q is the **shortest** path. Contradiction. Same can be said about Q .

Greedy choice is optimal

Theorem D1: Let $G = (V, E, W)$ be a weighted graph with nonnegative weights. Let S be a subset of V and let s be a member of S . Assume that $d[y]$ is the shortest distance in G from s to y , for each y in S . Let z be the next vertex chosen to go into S . If edge (y, z) is chosen to minimise $d[y] + W(y, z)$ over all edges with one vertex in S and one vertex in $V - S$, then the path consisting of a shortest path from s to y followed by the edge (y, z) is the shortest path from s to z .

Proof:


We will show that there is no other path from s to z that is shorter.

Greedy choice is optimal

$d[x_i]$: shortest distance from s to x_i


$d[y]$: shortest distance from s to y

Consider one **arbitrary vertex z**



Greedy choice is optimal

Which vertex to select to include into S ?



$s \rightarrow^* y \rightarrow z$: shortest path from s to z

Greedy choice is optimal

Proof of Theorem D1 (continued)


P: $s \rightarrow y \rightarrow z$ (shortest path for z)

$$W(P) = d[y] + W(y, z)$$

P': $s \rightarrow y \rightarrow u \rightarrow \dots \rightarrow z$
(an alternative shortest path)

$$W(P') = d[y] + W(y, u) + \text{distance from } u \text{ to } z$$

Because $d[y] + W(y, u) \geq d[y] + W(y, z)$,
and distance from u to z is nonnegative,
therefore $W(P) \leq W(P')$.



Edge (y, z) is chosen to minimise $d[y] + W(y, z)$ over all edges with one vertex in S and one vertex in $V - S$

Greedy choice is optimal

Proof of Theorem D1 (continued)

Let P be a shortest path from s to y followed by edge (y, z)

Let $W(P)$ = the distance travelled along P

Let P' = any shortest path different from P , i.e., $P' = s, z_1, \dots, z_k, \dots, z$


Assume that z_k is the first vertex in P' not in set S .

$$W(P) = d[y] + W(y, z)$$

$$W(P') = d[z_{k-1}] + W(z_{k-1}, z_k) + \text{distance from } z_k \text{ to } z$$

Note that: $d[z_{k-1}] + W(z_{k-1}, z_k) \geq d[y] + W(y, z)$

Since **distance from z_k to z** is non-negative, therefore, **$W(P) \leq W(P')$** .



Theorem D2 and Proof

Theorem D2: Given a directed weighted graph G with nonnegative weights and a source vertex s , Dijkstra's algorithm computes the shortest distance from s to each vertex of G that is reachable from s .

Proof (By induction):

We will show by induction that as each vertex v is added into set S , $d[v]$ is the shortest distance from s to v .

Basis:

The algorithm assigns $d[s]$ to zero when the source vertex s is added to S . So $d[s]$ is the shortest distance from s to s when S has the first vertex in it.


Theorem D2 and Proof (Continued)

Inductive Hypothesis:

Assume the theorem is true when S has k vertices. That is, assume $v_0, v_1, v_2, \dots, v_{k-1}$ are added where $d[v_1], d[v_2], \dots$ are the shortest distances.

When v_k is chosen by Dijkstra's algorithm, it means an edge (v_i, v_k) , where $i \in \{0, 1, 2, \dots, k-1\}$, is chosen to minimise $d[v_i] + W(v_i, v_k)$ among all edges with one vertex in S and one vertex not in S .





By Theorem D1, $d[v_k]$ is the shortest distance from source to v_k . So the theorem is true when S has $k + 1$ vertices.



Minimum Spanning Tree

Minimum Spanning Tree

Definition of Subgraph



A subgraph of a graph $G = (V, E)$ is a graph $G' = (V', E')$ such that $V' \subseteq V$ and $E' \subseteq E$ and $E' \subseteq V' \times V'$

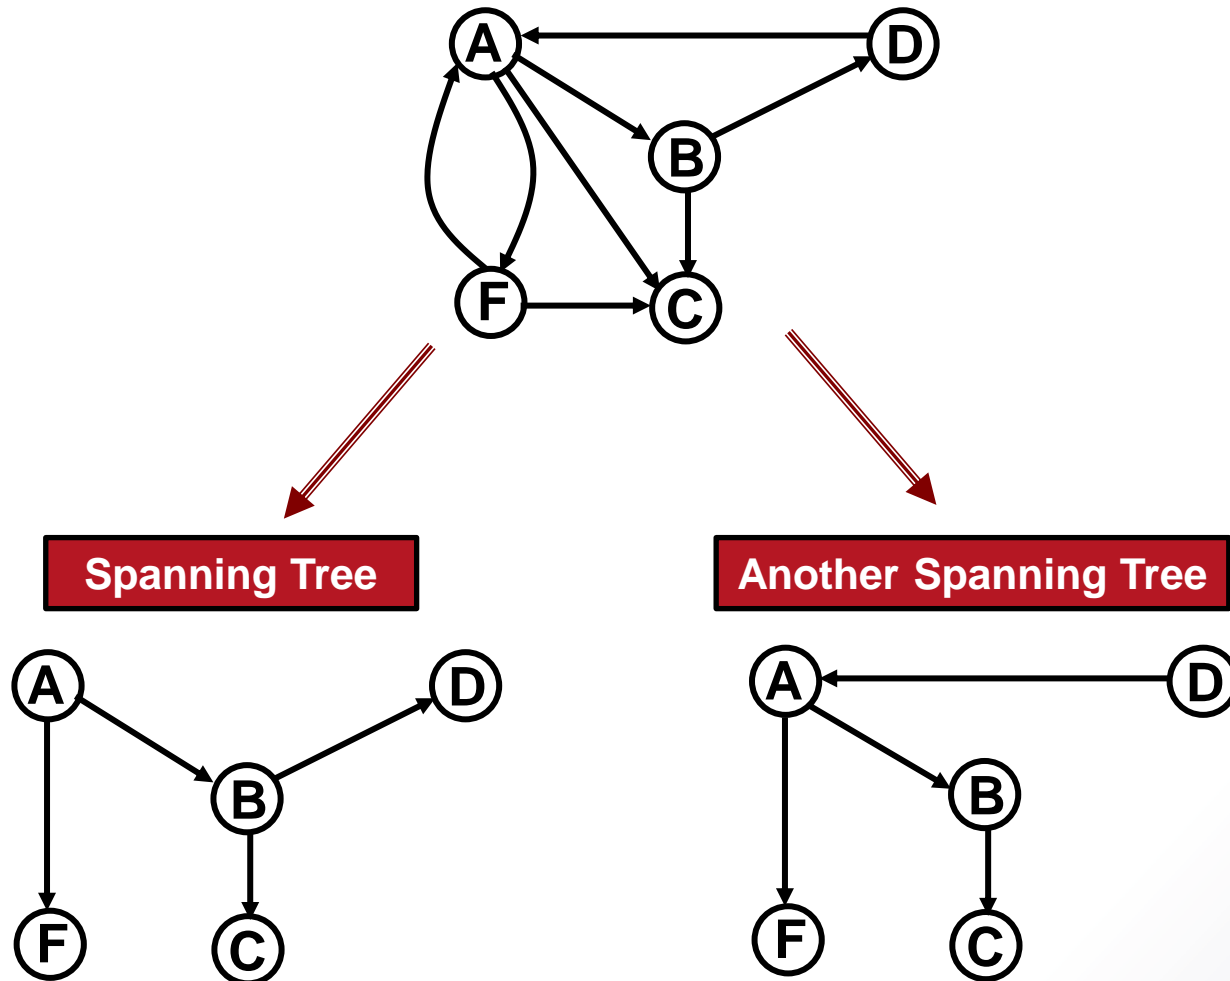
Definition of Spanning Tree

A connected, acyclic subgraph containing all the vertices of a graph.



Definition of Minimum Spanning Tree

A minimum-weight spanning tree in a weighted graph.


Spanning Tree




Minimum Spanning Tree



Spanning
Tree



Minimum
Spanning Tree




Minimum
Spanning Tree

Main Idea of Prim's Algorithm


Prim's Algorithm

- It works on undirected graph.
- It builds upon a single partial minimum spanning tree, at each step adding an edge connecting the vertex nearest to but not already in the current partial minimum spanning tree.
- At first a vertex is chosen, this vertex will be the first node in T .
- Set P is initialised: P = set of vertices not in tree T but are adjacent to some vertices in T .



Main Idea of Prim's Algorithm


Prim's Algorithm (Cont.)



- In every iteration in the Prim's Algorithm, a new vertex **u** from set **P** will be connected to the tree **T**. The vertex **u** will be deleted from the set **P**. The vertices adjacent to **u** and not already in **P** will be added to **P**.
- When all vertices are connected into **T**, **P** will be empty. This means the end of the algorithm.
- The new vertex in every iteration will be chosen by using **greedy** method, i.e. among all vertices in **P** which are connected to some vertices already inserted in the tree **T** but themselves are not in **T**, we choose one with the **minimum** cost.

An Example of Prim's Algorithm

Prim's MST



Black vertices: unseen vertices

Pink vertices: tree vertices

Blue vertices: fringe vertices

3 subsets of vertices


Prim's Algorithm classifies vertices into three disjoint categories:

- **Tree vertices** - in the tree being constructed so far
- **Fringe vertices** - not in the tree but adjacent to some vertices in the tree
- **Unseen vertices** - all others

Greedy choice of Prim's Algo

- Key step in the algorithm is the selection of a vertex from the fringe (which, of course, depends on the weights on incident edges).
- Prim's Algorithm always chooses a minimum weight edge from **tree** vertex to **fringe** vertex.

Main Idea of Prim's Algorithm



Choose $\min(w, x, y, z)$

Pseudocode of Prim's Algo

```
primMST(G, s, n) // outline of Prim's algorithm
{
    Initialise all vertices as unseen.
    Reclassify s as tree vertex.
    Reclassify all vertices adjacent to s as fringe.
    While (there are fringe vertices)
    {
        Select an edge of minimum weight between a tree
            vertex t and a fringe vertex v;
        Reclassify v as tree; add edge tv to the tree;
        Reclassify all unseen vertices adjacent to v as fringe.
    }
}
```


Implementing Prim's Algo

Data Structures Used:

- Array **d**: distance of a fringe vertex from the tree
- Array **pi**: vertex connecting a fringe vertex to a tree vertex
- Array **S**: whether a vertex is in the minimum spanning tree being built
- Priority queue **pq**: queue of fringe vertices in the order of the distances from the tree

At the end of the algorithm, array **pi** has the minimum spanning tree.

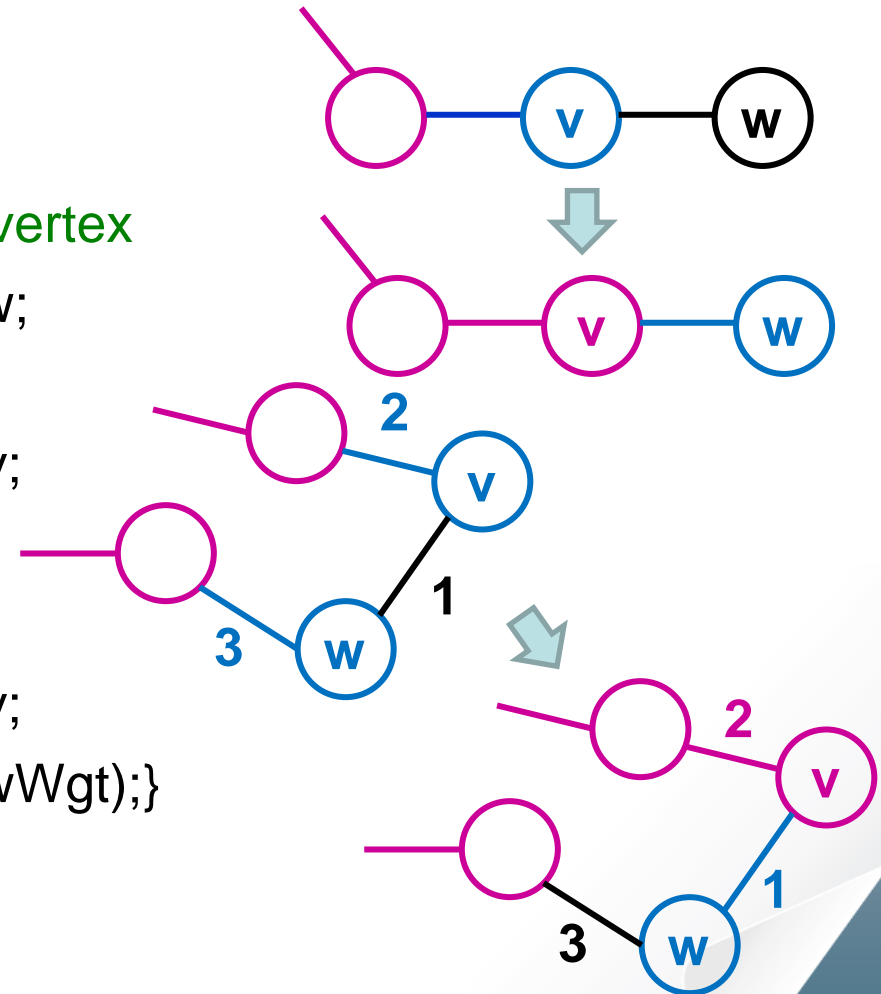
Implementing Prim's Algo

```
primMST(G, s, n) {  
    initialise priority queue pq as empty;  
    for each vertex v {  
        d[v] = infinity; S[v] = 0;  
        pi[v] = null pointer; }  
    d[s] = 0; S[s] = 1;  
    insert(pq, s, 0);  
    while (pq is not empty) {  
        u = getMin(pq); deleteMin(pq);  
        S[u] = 1;  
        updateFringe(pq, G, u); }  
}
```

Update Fringe Set of Vertices

```




updateFringe(pq, G, v) {
  for all vertices w adjacent to v {
    if (S[w] != 1) { //if w is not a tree vertex
      newWgt = weight of edge vw;
      if (d[w] == infinity) {
        d[w] = newWgt; pi[w] = v;
        insert(pq, w, newWgt);
      } else if (newWgt < d[w]) {
        d[w] = newWgt; pi[w] = v;
        decreaseKey(pq, w, newWgt);}
    } // if w is not a tree vertex
  } // for all vertices
}
  
```



MST Property

Minimum Spanning Tree Property

Let T be a spanning tree of G , where $G = (V, E, W)$ is a connected, weighted graph. Suppose that for every edge (u, v) of G that is not in T , if (u, v) is added to T it creates a cycle such that (u, v) is a maximum-weight edge on that cycle. Then T has the **Minimum Spanning Tree Property** (or **MST Property**, in short).



Lemma 1 and Proof

Lemma 1: In a connected weighted graph $G = (V, E, W)$, if T_1 and T_2 are two spanning trees that have the MST property, then they have the same total weight.

Proof by induction on k , the number of edges in T_1 but not T_2 (there are also k edges in T_2 but not in T_1).

Basis:

$k = 0$; i.e. $T_1 = T_2$. Therefore, they have the same weight.

Proof of Lemma 1 (continued)

Inductive hypothesis: For $k > 0$, assume the lemma holds when there are j differing edges where $0 \leq j < k$.

Let uv be the minimum weight edge among the differing edges (assume uv is in T_2 but not T_1).



Look at unique path in T_1 from u to v .

Suppose it is made up of w_0, w_1, \dots, w_p where $w_0 = u, \dots, w_p = v$.

This path must contain some edge different from T_2 's.

Let $w_i w_{i+1}$ be this differing edge.

By MST property of T_1 , $w_i w_{i+1}$ cannot be $> uv$'s weight.



Proof of Lemma 1 (continued)

But since uv was chosen to be the minimum weight among differing edges, $w_i w_{i+1}$ cannot have weight less than uv .

Therefore, $W(w_i w_{i+1}) = W(uv)$.

Add uv to T_1 (creating a cycle). Remove $w_i w_{i+1}$ leaving tree T'_1 (which has the same weight as T_1).

But T'_1 and T_2 differ only on $k-1$ edges.

So by inductive hypothesis, T'_1 and T_2 have the same total weight. Therefore, T_1 and T_2 have same weight.


Theorem 1 and Proof

Theorem 1: In a connected weighted graph $G = (V, E, W)$, a tree T is a minimum spanning tree if and only if T has the MST property.

Proof (Only if): Assume T is an MST for graph G .

Suppose T does not satisfy the MST property, i.e. there is some edge uv that is not in T such that adding uv creates a cycle, in which some other edge xy has weight $W(xy) > W(uv)$.

Then, by removing xy and adding uv , we create a new spanning tree whose total weight is $< W(T)$; This contradicts the assumption that T is an MST.



Proof of Theorem 1 (continued)

Theorem 1: In a connected weighted graph $G = (V, E, W)$, a tree T is a minimum spanning tree if and only if T has the MST property.

Proof (Only if): Assume T is an MST for graph G .

(Cont.)

(If) Assume T has MST property.

If T_{\min} is an MST, then T_{\min} has MST property by the first half of the proof.

By Lemma 1, $W(T) = W(T_{\min})$, so T is also an MST.

Prim's Algorithm is Optimal

Lemma 2: Let $G = (V, E, W)$ be a connected weighted graph; Let T_k be the tree with k vertices constructed by Prim's Algorithm, for $k = 1, 2, \dots, n$; and let G_k be the subgraph induced by the vertices of T_k . Then T_k has the MST property in G_k . (**Proof is not required**)

Theorem 2: Prim's Algorithm outputs a minimum spanning tree.

Proof:

- From Lemma 2, T_n has the MST property.
- By Theorem 1, T_n is a minimum spanning tree.

Priority Queue for MST (Optional)

- Inserted by order of priority (not chronological, as in 'normal' queues – FIFO)
- Elements to be inserted have a 'key' - contains the priority; element with highest priority will be selected first. [priority can be largest value (e.g. if we're computing max profit) or smallest value (e.g. if we're interested in min cost)]
- Think of pq as a sequence of pairs: (id_1, w_1) , $(id_2, w_2), \dots, (id_k, w_k)$. The order is in increasing w_i and id is a unique identifier for an element

Methods of Priority Queue (Optional)

The Priority Queue consists of:

Create: Constructor to set up PQ

isEmpty; getMin; getPriority: Access functions

insert; deleteMin; decreaseKey: Manipulation procedures

Insert(pq, id, w): Inserts (id, w) into an existing pq - position depends on w

decreaseKey(pq, id, neww): Rearranges pq based on new wt of element id

getMin(pq): Returns id_1 ;

getPriority(pq): Returns weight of min element

Summary

- Greedy algorithm is a general strategy to solve optimization problems
- Dijkstra's algorithm finds single-source shortest paths in a weighted graph of nonnegative edge weights
- Prim's algorithm finds the minimum spanning trees in weighted graphs
- Both are greedy algorithms, and use priority queue